

## 1. Project Title

**Retrieval-Augmented Generation (RAG) System for Sanskrit Text**

**Intern Name:** Vishal Dhaloch

---

## 2. Objective

The goal of this project is to implement a **CPU-only RAG pipeline** for Sanskrit text that can:

1. Load Sanskrit documents (.docx / .txt)
2. Preprocess and normalize text
3. Retrieve relevant context for a query
4. Generate Sanskrit answers using a lightweight LLM (TinyLlama)

This pipeline demonstrates **modern NLP concepts** (RAG) using accessible CPU-only resources.

## 3. Background / Motivation

### What is RAG?

RAG (Retrieval-Augmented Generation) is a system that combines:


1. **Information Retrieval** – retrieves relevant documents or chunks from a corpus.
2. **Generative LLM** – generates answers using the retrieved context.

### Advantages:

- Reduces hallucination by grounding responses in real documents
  - Efficient for domain-specific knowledge like Sanskrit texts
  - Modular — retriever and generator can be swapped independently
- 

## 4. Project Architecture

### 4.1 High-Level Flow



|

v

Answer in Sanskrit

## 4.2 Components

### 4.2.1 Loader (loader.py)

- Reads .docx and .txt files
- Consolidates all text into a single raw corpus
- Ensures **UTF-8/Devanagari compatibility**

**Functions:**

```
load_docx(path) # reads docx
```

```
load_text_file(path) # reads txt
```

```
load_documents(folder_path) # loads all files
```

### 4.2.2 Preprocessing (preprocess.py)

- Normalizes Sanskrit using **Unicode NFC**
- Removes unwanted characters except Devanagari and punctuation
- Splits text into overlapping **chunks** for retrieval

**Functions:**

```
normalize_sanskrit(text)
```

```
chunk_text(text, chunk_size=300, overlap=50)
```

---

### 4.2.3 Retriever (retriever.py)

- Implements **BM25 algorithm**
- Tokenizes chunks using nltk.word\_tokenize
- Retrieves **top-k relevant chunks** for a query

**Example:**

```
retriever = BM25Retriever(chunks)
```

```
top_chunks = retriever.retrieve("धर्मः किम्?", k=2)
```

### 4.2.4 Generator (generator.py)

- Uses **TinyLlama GGUF model**
- Fully CPU-only, lightweight (1.1B parameters)
- Handles n\_ctx=1024 tokens

- Generates Sanskrit answer from retrieved context

**Prompt Example:**

उत्तरं केवलं संस्कृतेन लिखत।

सन्दर्भः:

[retrieved chunks]

प्रश्नः:

धर्मः किम्?

उत्तरः:

**Token Safety:**

- Context truncated to prevent exceeding n\_ctx
- Max tokens for output limited (e.g., 64)

---

#### 4.2.5 RAG Pipeline (rag\_pipeline.py)

- Integrates Loader → Preprocessing → Retriever → Generator
- Class SanskritRAG exposes method ask(query)
- Handles:
  - Document loading
  - Cleaning
  - Chunking
  - Retrieval
  - Generation

**Usage:**

```
rag = SanskritRAG()
```

```
answer = rag.ask("धर्मः किम्?")
```

---

## 5. Sample Documents

- Stored in data/Rag-docs.docx
- Contains stories like:

- मूर्खभृत्यस्य कथा
- धर्म, कर्म, नीतिशास्त्र references

#### Example text snippet:

गोवर्धनदासः शंखनादम् शर्कराम् आनयितुम् आदिशति।

शंखनादः शर्कराम् जीर्णे वस्त्रे स्थापयति।

मार्ग सर्वा शर्करा नश्यति।

---

## 6. Sample Run

### 6.1 Input 1

धर्मः किम्?

### 6.1 Output

एवं पूर्णपद्याति स्वज्ञापनीति।

अप्रसाधिकरणाचार्यानुविधानम् क

#### Explanation:

- Retrieval fetched context about Dharma
- TinyLlama generated Sanskrit answer based on context
- CPU-only generation worked successfully

---

### 6.2 Input 2

मूर्खभृत्यस्य

### 6.2 Output

अरे शंखनाद गच्छति शर्कराम्

जीर्णवस्त्रात् मार्ग

सवर्पि शर्कर

#### Explanation:

- Context fetched from “servant story”
- TinyLlama generated partial story
- Output is consistent with retrieved document

## 7. Technical Details

Component	Library / Tool	Notes
Loader	python-docx, os	Reads files
Preprocessing	re, unicodedata	Devanagari normalization
Chunking	Custom function	Overlapping chunks (size=300)
Retriever	rank_bm25, nltk	BM25 scoring
Generator	llama-cpp-python	TinyLlama 1.1B GGUF
Model	TinyLlama	CPU-only inference
Context Safety	Hard truncation	Avoid n_ctx overflow

---

## 8. Installation Instructions

### Step 1: Create Environment

```
conda create -n lang6 python=3.10
```

```
conda activate lang6
```

### Step 2: Install Dependencies

```
pip install -r requirements.txt
```

```
pip install py-cpuinfo
```

### Step 3: Place Model

- Copy `tinyllama-1.1b-chat-v1.0.Q4_K_M.gguf` to `code/models/`

### Step 4: Run App

```
cd code
```

```
python app.py
```

### Step 5: Ask Sanskrit Questions

- Type query → Enter
- Type exit → Quit

---

## 9. Observations & Limitations

- TinyLlama is **not Sanskrit-trained**, output may be simple or incomplete
  - CPU-only execution is slower than GPU
  - Retrieval works perfectly — model may need **larger or domain-specific LLM** for fluent answer
-

## 10. References

1. Hugging Face TinyLlama GGUF: [Link](#)
2. rank\_bm25 Python Package
3. llama-cpp-python
4. Python docx, nltk, unicodedata

## 11. Screenshots

- Screenshot 1: Asking question धर्मः किम्?

Ask (or type 'exit'): धर्मः किम्?  
Answer:  
एवं पूर्वपद्यते स्वज्ञापनात्।  
अग्रसत्यधीन्द्रियाचार्यानुवृद्धानम् ॥  
Ask (or type 'exit'):  
Ask (or type 'exit'): exit

- Screenshot 2: Asking question मूर्खभृत्यस्य?

Ask (or type 'exit'): मूर्खभृत्यस्य  
Answer:  
अरे शङ्कनाद  
गच्छते शरुरातम्  
तीर्त्तावस्त्रात् मारगे  
श्रवणपि शरुर